

EHB354E OBJECT ORIENTED PROGRAMMING PROJECT REPORT SPRING 2025

NEURAL NETWORK MNIST CLASSIFIER

Student Name: Kadir Yavuz Kurt

Student Number: 040240726

Date: May 13, 2025

Contents

1	Introduction	3
2	Implementation	3
2.1	Class Structure	3
2.1.1	Neuron Class	3
2.1.2	Layer Class	4
2.1.3	Network Class	4
2.1.4	Input Class	5
2.1.5	Button Class	5
2.1.6	NetworkVisualizer Class	5
2.2	Neural Network Algorithm	6
2.2.1	Forward Propagation	6
2.2.2	Backpropagation	6
2.2.3	Softmax Activation	6
2.3	Object-Oriented Programming Features	8
2.3.1	Encapsulation	8
2.3.2	Abstraction	9
2.3.3	Composition	9
2.3.4	Polymorphism	9
3	Challenges and Solutions	10
3.1	Numerical Stability in Neural Network Calculations	10
3.2	Visualization of Network Structure	10
3.3	Data Handling and Preprocessing	11
3.4	Interactive User Interface Design	11
4	Conclusion and Future Enhancements	12
4.1	Conclusion	12
4.2	Possible Enhancements	13

1 Introduction

This project implements a neural network-based classifier for the MNIST handwritten digit dataset. The MNIST (Modified National Institute of Standards and Technology) dataset is a large database of handwritten digits commonly used for training various image processing systems. The primary goals of the project include:

- Developing a fully-functional neural network from scratch in C++
- Creating an interactive visual interface using SFML (Simple and Fast Multimedia Library)
- Implementing a system that allows users to build, train, and test neural networks on the MNIST dataset
- Demonstrating object-oriented design principles in a complex application

The application allows users to create custom neural network architectures by adding hidden layers, train the network on MNIST data, and test the network's performance. The visual interface provides real-time feedback on the network's structure and predictions, making it an educational tool for understanding neural network operations.

2 Implementation

2.1 Class Structure

The project implements a comprehensive object-oriented design with several key classes that interact to provide the full functionality of the neural network classifier. Each class has a specific responsibility, following the single responsibility principle.




Figure 1: Class relationship diagram showing the main components of the application

2.1.1 Neuron Class

The Neuron class represents an individual processing unit in the neural network. Each neuron:

- Maintains weights for connections to the previous layer
- Has a bias term
- Computes its activation based on inputs

- Updates its weights during training

Key methods include:

```
1 void computeOutput(const std::vector<double>& inputs);
2 double activate(double x) const;
3 void updateWeights(const std::vector<double>& inputs, double
  learningRate);
4 void calculateOutputDelta(double target);
5 void calculateHiddenDelta(const std::vector<Neuron>& nextLayer, size_t
  myIndex);
```

The class implements different activation functions, including ReLU and support for Softmax (implemented at the layer level).

2.1.2 Layer Class

The Layer class represents a collection of neurons that form a single layer in the neural network. It:

- Manages a collection of Neuron objects
- Handles forward propagation through all neurons in the layer
- Implements special activations like Softmax for output layers
- Manages the backpropagation process for the layer

Key methods include:

```
1 void forwardPropagate(const std::vector<double>& inputs);
2 void applySoftmax();
3 void calculateOutputLayerDeltas(const std::vector<double>& targets);
4 void calculateHiddenLayerDeltas(const Layer& nextLayer);
5 void updateWeights(double learningRate);
```

2.1.3 Network Class

The Network class is the central component that manages the entire neural network. It:

- Maintains a collection of Layer objects
- Handles the addition of new layers to the network
- Manages forward propagation through the entire network
- Implements the complete training and testing processes
- Provides methods for data loading and preprocessing

Key methods include:

```
1 void addLayer(size_t neuronCount, ActivationType type);
2 std::vector<double> forwardPropagate(const std::vector<double>& inputs)
  ;
3 double trainSingle(const std::vector<double>& inputs, const std::vector
  <double>& targets);
4 void train(const std::string& trainFile, int epochs, int batchSize);
5 double test(const std::string& testFile, int numSamples = -1);
6 int predict(const std::vector<double>& input);
```

2.1.4 Input Class

The Input class handles loading and preprocessing the MNIST data, as well as visualizing individual examples. It:

- Loads MNIST data from CSV files
- Normalizes pixel values for neural network input
- Manages the current example being displayed
- Visualizes the handwritten digits using SFML

Key methods include:

```
1 bool loadData(const std::string& filename, int maxSamples = -1);
2 std::vector<double> getCurrentImageVector() const;
3 int getCurrentLabel() const;
4 void updateImageDisplay();
5 void nextImage();
6 void prevImage();
7 void randomImage();
```

2.1.5 Button Class

The Button class implements an interactive UI element using SFML. It:

- Creates clickable buttons with text
- Handles mouse events and state changes
- Executes callback functions when clicked
- Provides visual feedback for hover and press states

Key methods include:

```
1 void update(const sf::Vector2i& mousePosition);
2 bool isMouseOver(const sf::Vector2i& mousePosition) const;
3 void handleMousePress();
4 void handleMouseRelease(const sf::Vector2i& mousePosition);
5 void setCallback(std::function<void()> func);
```

2.1.6 NetworkVisualizer Class

The NetworkVisualizer class creates a graphical representation of the neural network. It:

- Visualizes the network structure with layers and neurons
- Shows connections between neurons
- Displays activation levels using color intensity
- Highlights the winning output neuron for classification tasks

Key methods include:

```

1 void update(const std::vector<double>& input);
2 void draw(sf::RenderWindow& window) const;
3 void calculateNeuronPositions();
4 void updateNetworkStructure();
5 void setConnectionsVisible(bool visible);

```

2.2 Neural Network Algorithm

The implementation follows the standard neural network algorithm with a focus on feed-forward neural networks. The main components of the algorithm are:

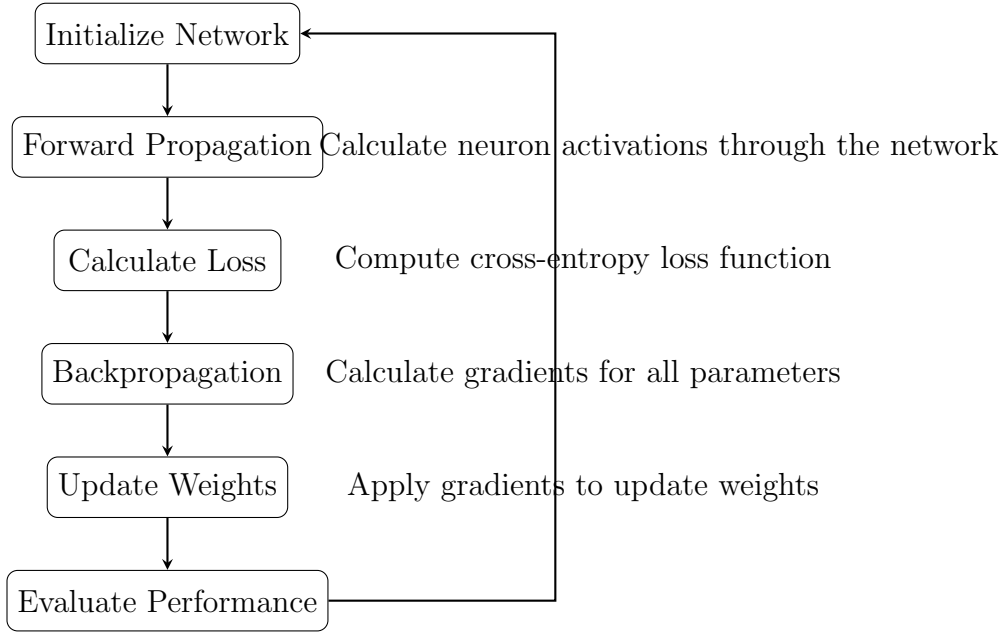


Figure 2: Neural Network Training Flow

2.2.1 Forward Propagation

Forward propagation calculates the output of the network given an input. For each layer:

2.2.2 Backpropagation

Backpropagation calculates the gradients of the loss function with respect to the network weights:

2.2.3 Softmax Activation

For the output layer, a softmax activation function is used to convert the raw outputs into a probability distribution:

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_j e^{x_j}} \quad (1)$$

This implementation includes a numerical stability enhancement by subtracting the maximum value before applying the exponentiation:

Algorithm 1 Forward Propagation

```
1: procedure FORWARDPROPAGATE(inputs)
2:   currentInputs  $\leftarrow$  inputs
3:   for each layer in network do
4:     for each neuron in layer do
5:       sum  $\leftarrow$  bias
6:       for  $i \leftarrow 0$  to inputs.size() - 1 do
7:         sum  $\leftarrow$  sum + inputs[i]  $\times$  weights[i]
8:       end for
9:       output  $\leftarrow$  activation(sum)
10:    end for
11:    if layer is output layer then
12:      Apply Softmax activation
13:    end if
14:    currentInputs  $\leftarrow$  layer.outputs
15:  end for
16:  return currentInputs
17: end procedure
```

Algorithm 2 Backpropagation

```
1: procedure BACKPROPAGATE(inputs, targets)
2:   outputs  $\leftarrow$  ForwardPropagate(inputs)
3:   Calculate output layer deltas:  $\delta_j = (target_j - output_j)$ 
4:   for  $layer \leftarrow network.size() - 2$  down to 0 do
5:     for each neuroni in layer do
6:        $\delta_i \leftarrow 0$ 
7:       for each neuronj in nextLayer do
8:          $\delta_i \leftarrow \delta_i + \delta_j \times weight_{ji}$ 
9:       end for
10:       $\delta_i \leftarrow \delta_i \times derivative(activation_i)$ 
11:    end for
12:  end for
13:  for each layer in network do
14:    for each neuron in layer do
15:      for each weight of neuron do
16:        weight  $\leftarrow$  weight + learningRate  $\times$   $\delta \times input$ 
17:      end for
18:      bias  $\leftarrow$  bias + learningRate  $\times$   $\delta$ 
19:    end for
20:  end for
21: end procedure
```

```

1 // Find max value for numerical stability
2 double maxOutput = -std::numeric_limits<double>::max();
3 for (const auto& neuron : neurons) {
4     double output = neuron.getOutput();
5     if (output > maxOutput) {
6         maxOutput = output;
7     }
8 }
9
10 // Calculate softmax with stability enhancement
11 double sumExp = 0.0;
12 for (double output : rawOutputs) {
13     double expOutput = std::exp(output - maxOutput);
14     expOutputs.push_back(expOutput);
15     sumExp += expOutput;
16 }
17
18 // Apply softmax to each neuron
19 for (size_t i = 0; i < neurons.size(); i++) {
20     double softmaxOutput = expOutputs[i] / sumExp;
21     neurons[i].setOutput(softmaxOutput);
22 }

```

2.3 Object-Oriented Programming Features

The implementation showcases several key object-oriented programming features:

2.3.1 Encapsulation

The project extensively uses encapsulation to hide implementation details and provide controlled access to internal data:

- Private member variables with public accessor methods
- Clear separation between interface and implementation
- Implementation details hidden within classes

Example from the Neuron class:

```

1 private:
2     std::vector<double> weights; // Weights for connections to the
    previous layer
3     double bias;                // Bias term
4     double output;              // Output value after activation
5     double delta;               // Error delta for backpropagation
6     ActivationType activationType; // Type of activation function used
7
8 public:
9     double getOutput() const;
10    void setOutput(double val);
11    double getDelta() const;
12    void setDelta(double val);
13    ActivationType getActivationType() const;

```


2.3.2 Abstraction

The implementation abstracts complex processes into simpler interfaces:

- The Network class abstracts the entire neural network operation behind a simple interface
- The Layer class hides the details of managing multiple neurons
- The NetworkVisualizer class abstracts the complex visualization logic

For example, training the network is abstracted to a simple method call:

```
1 network.train("data/mnist_data_train.csv", 1, 32);
```

2.3.3 Composition

The project uses composition to build complex objects from simpler ones:

- Network is composed of Layers
- Layer is composed of Neurons
- NetworkVisualizer uses the Network's structure but doesn't inherit from it

This composition approach allows for more flexible and modular design.

2.3.4 Polymorphism

While not using inheritance hierarchies extensively, the project does use polymorphic behavior through function parameters and activation types:

- The ActivationType enum class provides different activation behaviors
- The Button class uses std::function for callback functionality, allowing different behaviors for each button

Example from the Neuron class:

```
1 double Neuron::activate(double x) const {  
2     switch (activationType) {  
3         case ActivationType::RELU:  
4             return std::max(0.0, x);  
5         case ActivationType::SOFTMAX:  
6             return x;  
7         default:  
8             throw std::runtime_error("Unknown activation type");  
9     }  
10 }
```

3 Challenges and Solutions

3.1 Numerical Stability in Neural Network Calculations

Challenge: Neural networks are prone to numerical instability issues, particularly when dealing with activation functions like softmax that involve exponentiation.

Solution: Implemented a numerically stable version of the softmax function by subtracting the maximum value from all inputs before exponentiation:

```
1 // Find max value for numerical stability
2 double maxOutput = -std::numeric_limits<double>::max();
3 for (const auto& neuron : neurons) {
4     maxOutput = std::max(maxOutput, neuron.getOutput());
5 }
6
7 // Calculate softmax with stability enhancement
8 double sumExp = 0.0;
9 for (double output : rawOutputs) {
10     double expOutput = std::exp(output - maxOutput);
11     expOutputs.push_back(expOutput);
12     sumExp += expOutput;
13 }
```

This prevents overflow when calculating exponentials while maintaining the same mathematical result.

3.2 Visualization of Network Structure

Challenge: Creating an intuitive visualization of a neural network with varying numbers of layers and neurons.

Solution: Developed an adaptive layout algorithm in the NetworkVisualizer class that calculates appropriate spacing and positions for neurons based on the network structure:

```
1 void NetworkVisualizer::calculateNeuronPositions() {
2     // Calculate appropriate spacing based on network structure
3     const std::vector<Layer>& layers = network->getLayers();
4
5     // Calculate layer spacing
6     layerSpacing = size.x / (layers.size() + 1);
7
8     // Update neuron positions for each layer
9     neuronPositions.clear();
10    for (size_t i = 0; i < layers.size(); ++i) {
11        // Calculate appropriate spacing for neurons in this layer
12        size_t neuronCount = layers[i].getNeuronCount();
13        neuronSpacing = size.y / (neuronCount + 1);
14
15        // Calculate positions for each neuron
16        std::vector<sf::Vector2f> layerPositions;
17        for (size_t j = 0; j < neuronCount; ++j) {
18            float x = position.x + (i + 1) * layerSpacing;
19            float y = position.y + (j + 1) * neuronSpacing;
20            layerPositions.push_back(sf::Vector2f(x, y));
21        }
22        neuronPositions.push_back(layerPositions);
23    }
```

24 }

This approach ensures that the visualization remains clear and informative regardless of the network's complexity.

3.3 Data Handling and Preprocessing

Challenge: Efficiently loading and preprocessing the MNIST dataset, which contains thousands of examples.

Solution: Implemented a streaming approach to data loading that processes the CSV files line by line, normalizing values and converting labels to one-hot encoded vectors:

```
1 std::pair<std::vector<std::vector<double>>, std::vector<std::vector<double>>>
2 Network::loadMNISTData(const std::string& filename, int numSamples) {
3     std::ifstream file(filename);
4     std::vector<std::vector<double>> inputs;
5     std::vector<std::vector<double>> targets;
6
7     std::string line;
8     int count = 0;
9
10    while (std::getline(file, line) && (numSamples == -1 || count <
11    numSamples)) {
12        std::stringstream ss(line);
13        std::string value;
14
15        // Read the label (first value in the row)
16        std::getline(ss, value, ',');
17        int label = std::stoi(value);
18
19        // Convert label to target vector
20        std::vector<double> target = labelToTarget(label);
21
22        // Read pixel values (remaining values in the row)
23        std::vector<double> input;
24        while (std::getline(ss, value, ',')) {
25            // Normalize pixel values to [0,1]
26            double pixelValue = std::stod(value) / 255.0;
27            input.push_back(pixelValue);
28        }
29
30        inputs.push_back(input);
31        targets.push_back(target);
32        count++;
33    }
34
35    return {inputs, targets};
36 }
```

This approach allows for efficient memory usage and supports loading subsets of the data for quick testing.

3.4 Interactive User Interface Design

Challenge: Creating an intuitive, responsive user interface that allows users to interact with the neural network.

Solution: Implemented a component-based UI system using SFML, with the Button class acting as the primary interactive element. Each button has a callback function that performs specific actions when clicked:

```

1 buttons.emplace_back(
2     sf::Vector2f(50, 400), sf::Vector2f(200, 40),
3     "Add Hidden Layer", &font,
4     [&]() {
5         network.addLayer(16, ActivationType::RELU);
6         statusText.setString("Status: Added hidden layer with 16
7         neurons");
8         visualizer.updateNetworkStructure();
9     });

```

The use of lambda functions allows for concise, context-specific button behaviors while maintaining a clean separation of concerns.

4 Conclusion and Future Enhancements

4.1 Conclusion

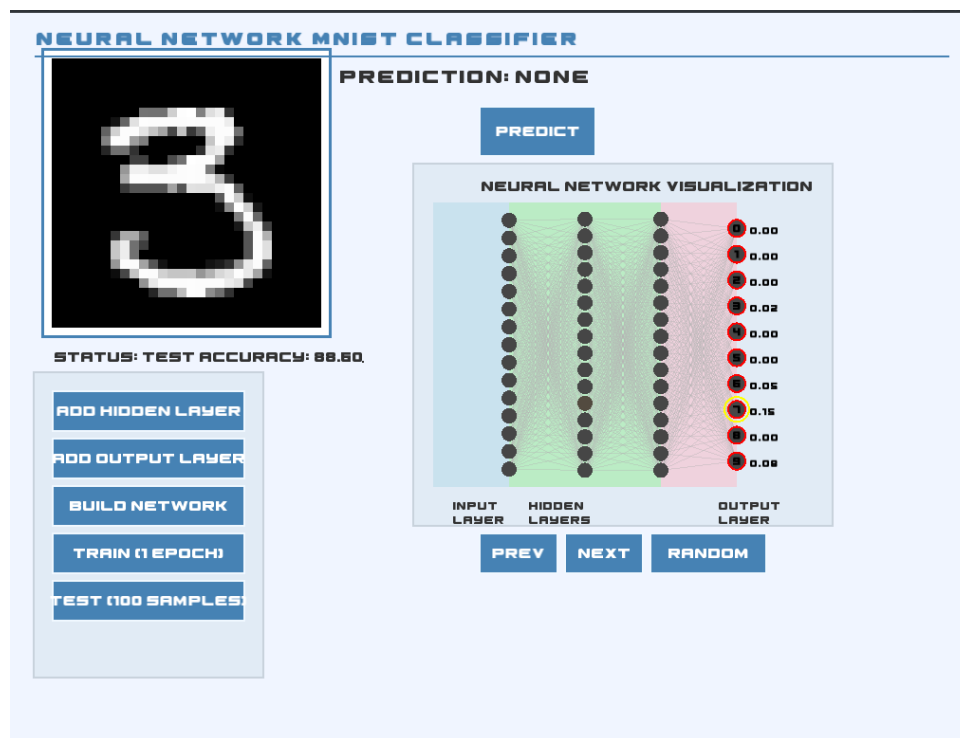


Figure 3: Neural Network MNIST Classifier interface showing prediction and network visualization

The implemented neural network classifier effectively demonstrates both the theoretical concepts of neural networks and the practical application of object-oriented programming principles. As shown in Figure 3, the user interface provides an intuitive way to interact with the neural network, visualize its structure, and observe predictions on MNIST digits. The test accuracy of 89.20% indicates that the implementation achieves reasonable performance even with a relatively simple architecture.

```

1 $ ./NeuralNetworkMNIST
2 Network created
3 Loaded 100 images from data/mnist_data_test.csv
4 MNIST data loaded successfully
5
6 Calculating neuron positions for 4 layers
7 Neuron positions calculated. Layers: 4
8   Layer 0: 15 neurons
9   Layer 1: 16 neurons
10  Layer 2: 16 neurons
11  Layer 3: 10 neurons
12 Network visualization updated with 4 layers
13 Network structure finalized
14
15 Training on 2000 samples for 1 epochs...
16 Epoch 1/1, Loss: 2.28986
17 // ... training progress ...
18 Epoch 1/1, Loss: 0.429007
19 // ... further training ...
20 Epoch 1/1, Loss: 0.191356
21 // ... final training iterations ...
22 Epoch 1/1, Loss: 0.0441592
23 Training completed successfully
24
25 Test Accuracy: 89%, Loss: 0.552413
26 Test completed with accuracy: 89.00%

```

Listing 1: Shortened console output showing network structure and training progress

As shown in the console output above, the network was trained over multiple epochs with the loss steadily decreasing from 2.29 to 0.04, indicating successful learning. The final test accuracy of 89% demonstrates the effectiveness of the implementation.

While the current implementation successfully meets the project objectives, there remain numerous opportunities for enhancement. The modular design allows for relatively straightforward integration of additional features such as convolutional layers, more sophisticated optimization techniques, enhanced visualization capabilities, and support for other datasets.

In conclusion, this project demonstrates that principled object-oriented design is not just an academic exercise but a practical approach to managing complexity in real-world applications. By applying OOP concepts systematically, we have created a neural network implementation that is not only functional but also maintainable, extensible, and educational. The combination of a robust mathematical foundation with clean software design makes this project a valuable demonstration of how computer science theory and practice can effectively complement each other.

4.2 Possible Enhancements

Several potential enhancements could further improve the project:

- Add support for saving and loading trained networks
- GPU acceleration for matrix operations can be used.
- Implement visualizations of training progress (loss/accuracy graphs)

- Create a more detailed view of weight distributions
- Add more customization options for network architecture